

# Tanmay Garg

CS20BTECH11063

## Deep Learning Assignment 1

```
In [ ]: import torch
import numpy as np
import matplotlib.pyplot as plt
import math
import torch.functional as F
import torch.nn as nn
```

### Q1

```
In [ ]: # Create a Linearly separable 2D dataset
def create_dataset(n=100, gamma=0.1):
    x = torch.randn(n, 2)
    y = torch.zeros(n)

    # for i in range(n):
    #     y[i] = 1 if x[i, 0] + gamma * x[i, 1] > 0 else -1
    # return x, y

    # Random initialize the weights and bias
    w = torch.randn(2)
    b = torch.randn(1)

    # Create the Dataset
    y = torch.sign(torch.matmul(x, w) + b)
    y[np.random.rand(n) < gamma] *= -1
    return x, y


# Plot the scatter plot with Legend
```

```
x, y = create_dataset(1000, 1)
scatter = plt.scatter(x[:, 0], x[:, 1], c=y, cmap=plt.cm.Paired, edgecolors='k')
xmin, xmax, ymin, ymax = plt.axis()
plt.title("Generated Dataset")
plt.xlabel("x1")
plt.ylabel("x2")
plt.legend(handles=scatter.legend_elements()[0], labels=["-1", "1"])
plt.show()
```

```
# print frequency percentage of labels
```

```
print("Percentage of -1 labels: ", (y == -1).sum().item() / len(y))
```

```
print("Percentage of 1 labels: ", (y == 1).sum().item() / len(y))
```



```
Percentage of -1 labels: 0.731
```

```
Percentage of 1 labels: 0.269
```

```
In [ ]: # split dataset into train and test
def split_dataset(x, y, train_ratio=0.8):
    n = len(x)
    train_size = int(train_ratio * n)
    x_train = x[:train_size]
    y_train = y[:train_size]
    x_test = x[train_size:]
    y_test = y[train_size:]
    return x_train, y_train, x_test, y_test

x_train, y_train, x_test, y_test = split_dataset(x, y, 0.8)
```

```
In [ ]: # Append 1 to x for bias
x_train = torch.cat((x_train, torch.ones(x_train.shape[0], 1)), dim=1)
print("Shape of Training x: ", x_train.shape)
print("Shape of Training y: ", y_train.shape)
# print(x)

Shape of Training x: torch.Size([800, 3])
Shape of Training y: torch.Size([800])
```

```
In [ ]: # Perceptron Training Algorithm
def perceptron_train(x, y, max_epochs=100):
    w = torch.zeros(3)
    k = 0
    for epoch in range(max_epochs):
        nb_changes = 0
        for i in range(x.size(0)):
            if x[i].dot(w) * y[i] <= 0:
                w = w + y[i] * x[i]
                nb_changes = nb_changes + 1
        if nb_changes == 0:
            # print('Stopping at Epoch: ', epoch)
            break
        k = k + 1
    # print('Number of changes: ', nb_changes)


    # return the weights and number of epochs
    return w, k
```

```
In [ ]: w, max_epochs_run = perceptron_train(x_train, y_train, 100)
print("W = ", w, " Max Epochs Run = ", max_epochs_run)
```

W = tensor([-16.7979, -9.3326, -12.0000]) Max Epochs Run = 19

```
In [ ]: # plot the decision boundary using test data
x1 = np.linspace(-5, 5, 100)
x2 = -(w[0] * x1 + w[2]) / w[1]
plt.plot(x1, x2)
scatter = plt.scatter(x_test[:, 0], x_test[:, 1], c=y_test, cmap=plt.cm.Paired, edgecolors='k')
plt.ylim(ymin, ymax)
plt.xlim(xmin, xmax)
plt.legend(handles=scatter.legend_elements()[0], labels=["-1", "1"])
plt.show()

# print accuracy on test data
x_test = torch.cat((x_test, torch.ones(x_test.shape[0], 1)), dim=1)
y_pred = torch.sign(x_test @ w)
print("Accuracy = ", (torch.sum(y_pred == y_test) / y_test.shape[0]).item() * 100, "%")
```



Accuracy = 99.50000047683716 %

```

In [ ]: # Running the perceptron training algorithm for different values of gamma for multiple trials
num_trials = 5
gamma_val = []
k_val = []
for gamma in np.linspace(0.1, 2, 20):
    acc = 0
    k_avg = 0
    for i in range(num_trials):
        x, y = create_dataset(1000, gamma)

        # split dataset into train and test
        x_train, y_train, x_test, y_test = split_dataset(x, y, 0.8)
        # x = x_train
        # y = y_train


        x_train = torch.cat((x_train, torch.ones(x_train.shape[0], 1)), dim=1)
        w, k = perceptron_train(x_train, y_train, 100)
        k_avg = k_avg + k

        x_test = torch.cat((x_test, torch.ones(x_test.shape[0], 1)), dim=1)
        y_pred = torch.sign(x_test @ w)
        acc = acc + torch.sum(y_pred == y_test) / y_test.shape[0]
        # print('-----')
    k_avg = k_avg / num_trials
    gamma_val.append(gamma)
    k_val.append(k_avg)
    print('Gamma: %f, Accuracy: %f, Avg. Epochs: %f' % (gamma, acc / num_trials, k_avg))
    # print('-----')

plt.plot(gamma_val, k_val)
plt.xlabel('Gamma')
plt.ylabel('Avg. Epochs')
plt.show()

```

Gamma: 0.100000, Accuracy: 0.821000, Avg. Epochs: 100.000000  
Gamma: 0.200000, Accuracy: 0.613000, Avg. Epochs: 100.000000  
Gamma: 0.300000, Accuracy: 0.525000, Avg. Epochs: 100.000000  
Gamma: 0.400000, Accuracy: 0.502000, Avg. Epochs: 100.000000  
Gamma: 0.500000, Accuracy: 0.505000, Avg. Epochs: 100.000000  
Gamma: 0.600000, Accuracy: 0.510000, Avg. Epochs: 100.000000  
Gamma: 0.700000, Accuracy: 0.595000, Avg. Epochs: 100.000000  
Gamma: 0.800000, Accuracy: 0.701000, Avg. Epochs: 100.000000  
Gamma: 0.900000, Accuracy: 0.779000, Avg. Epochs: 100.000000  
Gamma: 1.000000, Accuracy: 0.995000, Avg. Epochs: 32.800000  
Gamma: 1.100000, Accuracy: 0.999000, Avg. Epochs: 22.400000  
Gamma: 1.200000, Accuracy: 1.000000, Avg. Epochs: 44.600000  
Gamma: 1.300000, Accuracy: 0.996000, Avg. Epochs: 4.200000  
Gamma: 1.400000, Accuracy: 0.999000, Avg. Epochs: 36.400000  
Gamma: 1.500000, Accuracy: 0.999000, Avg. Epochs: 16.200000  
Gamma: 1.600000, Accuracy: 0.998000, Avg. Epochs: 30.600000  
Gamma: 1.700000, Accuracy: 0.998000, Avg. Epochs: 30.600000  
Gamma: 1.800000, Accuracy: 0.996000, Avg. Epochs: 30.600000  
Gamma: 1.900000, Accuracy: 0.995000, Avg. Epochs: 14.200000  
Gamma: 2.000000, Accuracy: 0.995000, Avg. Epochs: 50.600000




We can see that as the value of  $\gamma$  increases, the average number of epochs required to converge decreases. This is because as  $\gamma$  increases, the data becomes more and more linearly separable and hence the number of epochs required to converge decreases.

## Q2

We will be using Hinge loss function with scratch implementation of gradient descent algorithm

```
In [ ]: # Create a Linearly separable 2D dataset
x, y = create_dataset(1000, 1)
scatter = plt.scatter(x[:, 0], x[:, 1], c=y, cmap=plt.cm.Paired, edgecolors='k')
xmin, xmax, ymin, ymax = plt.axis()
plt.legend(handles=scatter.legend_elements()[0], labels=["-1", "1"])
plt.show()
```

```
print("Percentage of -1 labels: ", (y == -1).sum().item() / len(y))
print("Percentage of 1 labels: ", (y == 1).sum().item() / len(y))
```



```
Percentage of -1 labels: 0.718
Percentage of 1 labels: 0.282
```

```
In [ ]: x_train, y_train, x_test, y_test = split_dataset(x, y, 0.8)

# Append 1 to x for bias

x_train = torch.cat((x_train, torch.ones(x_train.shape[0], 1)), dim=1)
print("Shape of Training x: ", x_train.shape)
print("Shape of Training y: ", y_train.shape)

Shape of Training x: torch.Size([800, 3])
Shape of Training y: torch.Size([800])
```

```
In [ ]: # Gradient Descent Algorithm for Hinge Loss
def hinge_gradient_descent(x, y, lr=0.1, max_epochs=100):
    w = torch.zeros(3)
```



```

k = 0
for epoch in range(max_epochs):
    nb_changes = 0
    for i in range(x.size(0)):
        if x[i].dot(w) * y[i] < 1:
            w = w + lr * (y[i]*x[i])
            nb_changes = nb_changes + 1
    if nb_changes == 0:
        # print('Stopping at Epoch: ', epoch)
        break
    k = k + 1
    # print('Number of changes: ', nb_changes)

    # return the weights and number of epochs
return w, k

```

```

In [ ]: w, max_epochs = hinge_gradient_descent(x_train, y_train, 0.1, 100)
print("W = ", w, " Max Epochs Run = ", max_epochs)

```

```

W = tensor([-17.2833,  6.0684, -10.1000]) Max Epochs Run = 100


```

```

In [ ]: # plot the decision boundary
x1 = np.linspace(-4, 4, 100)
x2 = -(w[0] * x1 + w[2]) / w[1]
plt.plot(x1, x2)
scatter = plt.scatter(x[:, 0], x[:, 1], c=y, cmap=plt.cm.Paired, edgecolors='k')
plt.xlim(xmin, xmax)
plt.ylim(ymin, ymax)
plt.legend(handles=scatter.legend_elements()[0], labels=["-1", "1"])
plt.show()

# print accuracy
x_test = torch.cat((x_test, torch.ones(x_test.shape[0], 1)), dim=1)
y_pred = torch.sign(x_test @ w)
print("Accuracy = ", (torch.sum(y_pred == y_test) / y_test.shape[0]).item() * 100, "%")

```



Accuracy = 99.00000095367432 %

## Q3

```
In [ ]: # create dataset with concentric circles
# import sklearn.datasets as skdata
from sklearn.datasets import make_circles

def create_concentric_dataset(n_samples, factor=0.9, noise=0.05):
    radius = np.random.rand(n_samples) * factor
    angle = np.random.rand(n_samples) * 2 * np.pi


    dataset = np.column_stack((radius * np.cos(angle), radius * np.sin(angle))) + np.random.normal(0, noise, (n_samples, 2))

    labels = np.zeros(n_samples)
    labels[radius > factor / 2] = 1
```

```
return dataset, labels
```

```
x, y = create_concentric_dataset(n_samples=1250, factor=0.9, noise=0.05)
scatter = plt.scatter(x[:, 0], x[:, 1], c=y, cmap=plt.cm.Paired, edgecolors='k')
plt.legend(handles=scatter.legend_elements()[0], labels=["0", "1"])
plt.show()

print("Percentage of 0 labels: ", (y == 0).sum().item() / len(y))
print("Percentage of 1 labels: ", (y == 1).sum().item() / len(y))
```



```
Percentage of 0 labels: 0.4872
Percentage of 1 labels: 0.5128
```

```
In [ ]: print(x.shape)
        print(y.shape)
        x = torch.from_numpy(x).float()
        y = torch.from_numpy(y).float()

        x_train, y_train, x_test, y_test = split_dataset(x, y, 0.8)
```

```
print("Shape of Training x: ", x_train.shape)
print("Shape of Training y: ", y_train.shape)
```

```
(1250, 2)
```

```
(1250,)
```

```
Shape of Training x: torch.Size([1000, 2])
```

```
Shape of Training y: torch.Size([1000])
```

```
In [ ]: # Create MLP with 1 hidden layer from scratch
```

```
class MLP:
    def __init__(self, x, y, hidden_size=4, lr=0.1) -> None:
        self.x = x
        self.y = y
        self.input_size = x.shape[1]
        self.hidden_size = hidden_size
        self.output_size = 1
        self.lr = lr

        # Weights and Biases
        self.w1 = torch.randn(self.input_size, self.hidden_size)
        # print("w1: ", self.w1.shape)
        self.b1 = torch.randn(1) * torch.randn(self.hidden_size)
        # print("b1: ", self.b1.shape)
        self.w2 = torch.randn(self.hidden_size, self.output_size)
        # print("w2: ", self.w2.shape)
        self.b2 = torch.randn(1)

    # Sigmoid Activation Function
    def sigmoid(self, x):
        return 1 / (1 + torch.exp(-x))

    # Sigmoid Derivative
    def sigmoid_derivative(self, x):
        return x * (1 - x)

    # Hinge Loss Function
    def hinge_loss(self, y_pred, y):
        # print(y_pred.shape)
        # print(y.shape)
        # print((y_pred * y).shape)
        return torch.max(torch.zeros_like(y_pred), 1 - y_pred * y)

    # Square Loss Function
```

```

def square_loss(self, y_pred, y):
    return (y_pred - y) ** 2

# Hinge Loss Derivative
def hinge_loss_derivative(self, y_pred, y):
    # print(y_pred.shape)
    # print(y.shape)
    return -y * (y * y_pred < 1)

# Square Loss Derivative
def square_loss_derivative(self, y_pred, y):
    return 2 * (y_pred - y)

# Binary Cross Entropy Loss Function
def binary_cross_entropy(self, y_pred, y):
    return -y * torch.log(y_pred) - (1 - y) * torch.log(1 - y_pred)

# Binary Cross Entropy Loss Derivative
def binary_cross_entropy_derivative(self, y_pred, y):
    return -(y / y_pred) + ((1 - y) / (1 - y_pred))

# Forward Propagation
def forward(self, x):
    self.z1 = x @ self.w1 + torch.Tensor.repeat(self.b1, x.shape[0], 1)
    # print("z1: ", self.z1.shape)

    self.a1 = self.sigmoid(self.z1)
    # print("a1: ", self.a1.shape)

    self.z2 = self.a1 @ self.w2 + self.b2
    # print("z2: ", self.z2)

    self.a2 = self.sigmoid(self.z2)
    # print("a2: ", self.a2)

    # print("w2: ", self.w2)
    # print("w1: ", self.w1)
    return self.a2

# Backward Propagation
def backward(self, x, y):
    y = y.reshape(-1, 1)
    # print("x shape: ", x.shape)

```

```

self.loss_a2 = self.binary_cross_entropy(self.a2, y)
# print("Loss_a2 shape: ", self.loss_a2.shape)

# Calculate Gradients
self.dL_d_a2 = self.binary_cross_entropy_derivative(self.a2, y)
# print("dL_d_a2 shape: ", self.dL_d_a2.shape)

self.da2_dz2 = self.sigmoid_derivative(self.a2)
# print("da2_dz2 shape: ", self.da2_dz2.shape)

self.dz2_d_w2 = self.a1.T
# print("dz2_d_w2 shape: ", self.dz2_d_w2.shape)

self.dz2_d_b2 = torch.ones_like(self.z2)
# print("dz2_d_b2 shape: ", self.dz2_d_b2.shape)

self.dL_d_w2 = self.dz2_d_w2 @ (self.dL_d_a2 * self.da2_dz2)
# print("dL_d_w2 shape: ", self.dL_d_w2.shape)

# self.dL_d_b2 = (self.dz2_d_b2 * (self.dL_d_a2.T @ self.da2_dz2)).reshape(-1)
self.dL_d_b2 = ((self.dL_d_a2 * self.da2_dz2).T @ self.dz2_d_b2).reshape(-1)
# print("dL_d_b2 shape: ", self.dL_d_b2.shape)

self.da1_dz1 = self.sigmoid_derivative(self.a1)
# print("da1_dz1 shape: ", self.da1_dz1.shape)

self.dz1_dw1 = x
# print("dz1_dw1 shape: ", self.dz1_dw1.shape)

self.dz1_x = self.w1

self.dz2_d_a1 = self.w2

# print("dz2_d_a1 shape: ", self.dz2_d_a1.shape)
# print("dz1_x shape: ", self.dz1_x.shape)
# self.dz1_d_b1 = torch.ones_like(self.a1)
# print("dz1_d_b1 shape: ", self.dz1_d_b1.shape)

# print("ter: ", ((self.dL_d_a2 * self.da2_dz2) @ self.w2.T).shape)
self.dL_dw1 = self.dz1_dw1.T @ (((self.dL_d_a2 * self.da2_dz2) @ self.w2.T) * self.da1_dz1)

# print("dL_dw1 shape: ", self.dL_dw1.shape)
# print("ter ", (((self.dL_d_a2 * self.da2_dz2) @ self.w2.T) @ self.da1_dz1.T).shape)
# print((((self.dL_d_a2 * self.da2_dz2).T @ (self.da1_dz1 * self.dz1_d_b1))).T.shape)

```

```

self.dL_db1 = (((self.dL_d_a2 * self.da2_dz2) @ self.w2.T) * self.da1_dz1).sum(axis=0)

# print("dL_db1 shape: ", self.dL_db1.shape)

# Updating Weights and Biases
self.w2 -= self.lr * self.dL_d_w2
self.b2 -= self.lr * self.dL_d_b2
self.w1 -= self.lr * self.dL_dw1
self.b1 -= self.lr * self.dL_db1

# Training the Model
def train(self, epochs=100):
    loss = []
    for epoch in range(epochs):
        self.forward(self.x)
        # print("W1 in forward: ", self.w1)
        # print("W2 in forward: ", self.w2)
        # print("B2 in forward: ", self.b2)
        # print("A2 in forward: ", self.a2)
        self.backward(self.x, self.y)
        # print("W1 in backward: ", self.w1)
        # print("W2 in backward: ", self.w2)
        # print("B2 in backward: ", self.b2)

        loss.append(self.loss_a2.mean())
        if epoch % 10 == 0:
            print("Epoch: ", epoch, " Loss: ", self.loss_a2.mean())
    # print("Final Predicted: ", self.forward(self.x))
    # print("Ground Truth: ", self.y)
    # print("z2: ", self.z2)

    # Plotting the Loss Curve
    plt.figure(figsize=(20, 5))
    plt.xticks(np.arange(0, epochs + 10, 10))
    plt.plot(loss)
    plt.xlim(-1, epochs)
    plt.ylim(min(loss), max(loss) + 0.1)
    plt.yticks(np.arange(0, max(loss) + 0.1, 0.1))
    plt.show()
    return

def predict(self, x, threshold=0.5):

```

```

        predicted = self.forward(x)
        # print(torch.round(predicted))
        # apply threshold of 0.6 for label
        # print(predicted)

        return predicted > threshold

    def accuracy(self, x, y):
        # print(self.predict(x).shape)
        y = y.reshape(-1, 1)
        return torch.sum(self.predict(x) == y) / y.shape[0]

```

```

In [ ]: # create MLP
mlp = MLP(x_train, y_train, hidden_size=4, lr=0.01)
mlp.train(epochs=200)

print("Accuracy on Test set: ", mlp.accuracy(x_test, y_test).item() * 100, "%")

# plot the predicted labels on the test set
scatetr = plt.scatter(x_test[:, 0], x_test[:, 1], c=mlp.predict(x_test).detach().numpy(), cmap=plt.cm.Paired, edgecolors='k')
plt.legend(handles=scatetr.legend_elements()[0], labels=['0', '1'])
plt.show()


```

```

Epoch: 0   Loss: tensor(1.4430)
Epoch: 10  Loss: tensor(1.0920)
Epoch: 20  Loss: tensor(1.0762)
Epoch: 30  Loss: tensor(1.0729)
Epoch: 40  Loss: tensor(1.0751)
Epoch: 50  Loss: tensor(1.0746)
Epoch: 60  Loss: tensor(1.0250)
Epoch: 70  Loss: tensor(0.8837)
Epoch: 80  Loss: tensor(0.7067)
Epoch: 90  Loss: tensor(0.6489)
Epoch: 100 Loss: tensor(0.5653)
Epoch: 110 Loss: tensor(0.3599)
Epoch: 120 Loss: tensor(0.2105)
Epoch: 130 Loss: tensor(0.1802)
Epoch: 140 Loss: tensor(0.1703)
Epoch: 150 Loss: tensor(0.1643)
Epoch: 160 Loss: tensor(0.1599)
Epoch: 170 Loss: tensor(0.1563)
Epoch: 180 Loss: tensor(0.1533)
Epoch: 190 Loss: tensor(0.1506)

```





Accuracy on Test set: 96.39999866485596 %

